




Computing with neurons



Session 1: Neural Engineering







Centre for Theoretical Neuroscience University of Waterloo

- Chris Eliasmith





CTN and ABR





Understanding the mind

*What I cannot create,
I do not understand.*

- What are the algorithms underlying cognition?
 - What is the mind doing?
- How can we know what the mind is doing?
 - How can we test these theories?



Understanding the mind



- Build computer models of the mind
- Mechanistic models
 - Internal components that map onto real system
- Why do we need a computational model?
 - Not analytically tractable



Cognitive Modelling

- Choose a phenomenon
 - Examine human behaviour
 - Build a computer program
 - Compare behaviours of program and human
- Many domains
 - Memory, mental arithmetic, reward learning
- Problem
 - How do we know if we're right?



Cognitive Architectures

- More constrained approach
 - Define a bunch of basic modules
 - Declarative memory
 - Visual recognition
 - Hand movement
 - Procedural memory
 - Use the same set of modules to do many tasks
 - It's not like we suddenly get new brain areas for each new task



Cognitive Architectures

- ACT-R

- Declarative memory
- Procedural memory

- IF-THEN rules

- IF I'm counting and I'm at THREE then go to FOUR

- Parameter values

- $d = 0.5$

- 50 ms per rule

- Found by looking at human data across many conditions

- What are the limits on the procedural rules?

$$B_i = \ln\left(\sum_{j=1}^n t_j^{-d}\right) + \beta_i$$



Cognitive Architectures

- Wide variety of tasks
 - Mental arithmetic
 - Estimating time
 - Visual search
 - Air-traffic control
 - Military squad co-ordination
 - Language interpretation
 - Driving a car
 - Dialing a phone number
 - Driving a car while dialing a phone number



Cognitive Architectures

- How does this help?
 - More constraints
- Same components do many different tasks
- Parameter values shouldn't change (much)
 - Or theory can say when they change
- Predicting many different aspects of behaviour with a small set of components



What about the brain?

- Should we pay attention to it?
- Why would it matter for algorithms?
 - Why not just look for the best algorithm?
 - Why constrain ourselves?



Advantage 1

More predictions

- A brain-based model will predict more than just overt behaviour
 - Connectivity
 - Firing patterns
 - Results of lesions
 - Timing
 - Effects of drugs



Advantage 2

Different algorithms

- Infinite numbers of algorithms to consider
- We implement algorithms on computers
 - So we are biased toward considering algorithms *that are easy to program*
- Instead, let's determine the types of algorithms that neurons would be good at implementing
 - Then make software tools to make those types of algorithms easy to program



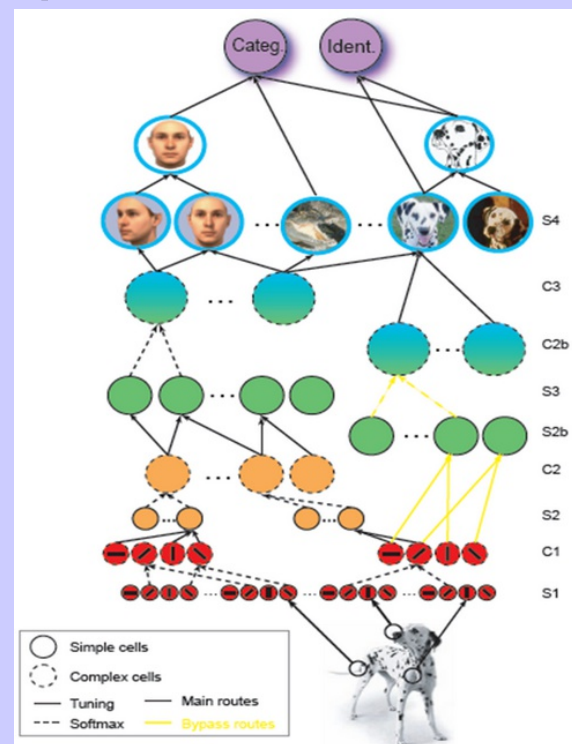

The Brain

- What is the brain?
- How should we think about the brain?
 - 140,000,000,000,000,000,000,000,000 atoms?
 - 100,000,000,000 neurons?
 - 52 brain areas?



Connectionism


- Neural networks
 - Many components
 - Many connections
 - Components add their inputs, perform some non-linearity to produce outputs





Connectionism


- How do we decide what the components can do?
 - Common choice: sigmoid neuron
 - Why that one?
 - Easy to program
- How do we get connection weights?
 - Start random, apply learning rule
 - Gets better and better at task (maybe)
 - Lots of computing needed





Neural Engineering Framework



- (Eliasmith & Anderson, 2003)
- Is there another way?
 - What are realistic neurons good at computing?
 - Can this help resolve the connection weight problem?





A neuron


- What is a neuron really like?





A neuron




- What is a neuron really like?





A neuron



- What is a neuron really like?



Link: [crumb of mouse brain](#)



A neuron



r_m : Membrane resistance


r_l : Longitudinal resistance

c_m : Capacitance due to electrostatic forces

$$\frac{r_m}{r_l} \frac{\partial^2 V}{\partial x^2} = c_m r_m \frac{\partial V}{\partial t} + V$$




Many neurons (and synapses)





Many many neurons (and synapses)



(Markram et al., Cell, 2015)

31,000 neurons; 37 million synapses



How much detail?

- Do we need all that complexity?
 - How do we know when to stop?
- How do other sciences deal with this?



How much detail?

- Do we need all that complexity?
 - How do we know when to stop?
- How do other sciences deal with this?
 - Physics (gravity): sometimes Newton is enough detail, sometimes you need Einstein
- The level of detail needed depends on the question being asked
 - e.g. drug effects may require a detailed model
 - But do we need it for understanding behaviour?



Behaving Systems

- Brains are for behaving
 - Sensory input, muscle outputs
- If we want computational neuroscience to explain what people do and how they do it, then the models need to produce that behaviour
 - Given similar inputs as the real system:
 - Produce similar outputs
- But behaviour requires many more neurons....



A neuron


- Could use a full supercomputer to simulate one neuron
- Have to make some abstraction
 - Start with something simple and uncontroversial
 - But everything we do could also be applied to more complex neurons

$$I(t) - \frac{V_m(t)}{R_m} = C_m \frac{dV_m(t)}{dt}$$



Tuning Curves


- How do neurons represent?
 - What is the relationship between the activity of a neuron a and the variable being represented x ?
 - Sometimes this is easy:





Tuning Curves

- Other times, not so much





Tuning Curves

- Let's break these tuning curves down into two aspects
 - Mapping from x (the variable) to J (current)
 - This is about how this neuron's inputs are organized
 - Mapping from J (current) to a (activity)
 - This is about the intrinsic response
 - This can be as complex as you want (assuming you have the compute power to do so)
 - a can be spikes or rates
 - I'm going to plot it as rates for now



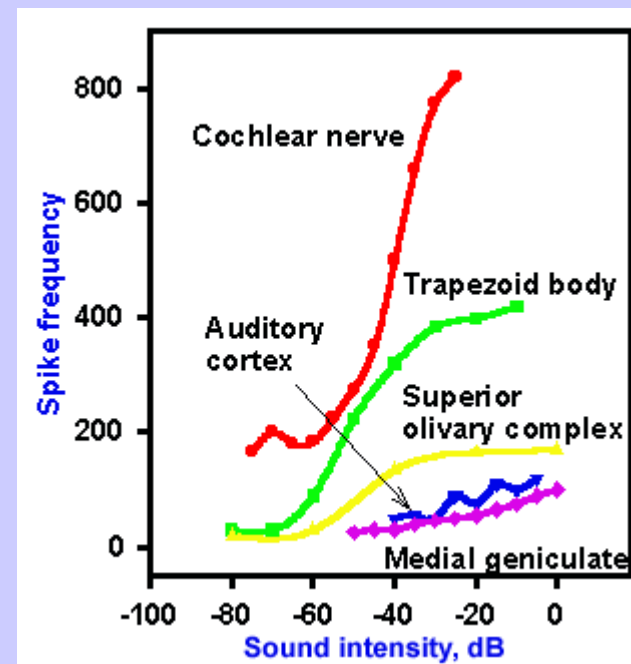

Tuning Curves

- Easy case:

- $J = \alpha x + J^{bias}$

- For activity, use standard LIF model for now

```
plot(x, n.rates(x, gain=1, bias=50), 'b') # x*1+50  
plot(x, n.rates(x, gain=0.1, bias=10), 'r') # x*0.1+10  
plot(x, n.rates(x, gain=0.5, bias=5), 'g') # x*0.05+5  
plot(x, n.rates(x, gain=0.1, bias=4), 'c') # x*0.1+4)
```








Tuning Curves

- What about these?
 - Mapping from x to J

$$J = -\alpha x + J^{bias}$$





- Okay, what about these?





Tuning Curves



- There's usually some x which gives a maximum firing rate
 - ...and thus a maximum J
- Firing rate (and J) decrease as you get farther from the preferred x value
 - So something like $J = \alpha[\text{sim}(x, x_{pref})] + J^{bias}$
- What sort of similarity measure?
- Let's think about x for a moment
 - x can be anything... scalar, vector, etc.
 - Does thinking of it as a vector help?



Tuning Curves


- Here is the general form we use for everything (it has both 'mappings' in it)
- $a_i = G_i[\alpha_i x \cdot e_i + J_i^{bias}]$
 - α is a gain term (constrained to always be positive)
 - J^{bias} is a constant bias term
 - e is the *encoder*, or the *preferred direction vector*
 - G is the neuron model
 - i indexes the neuron
- To simplify life, we always assume e is of unit length
 - Otherwise we could combine α and e
- In the 1D case, e is either +1 or -1
- In higher dimensions, what happens?



Tuning Curves

- 2-dimensional x

$$a_i = G_i[\alpha_i x \cdot e_i + J_i^{bias}]$$

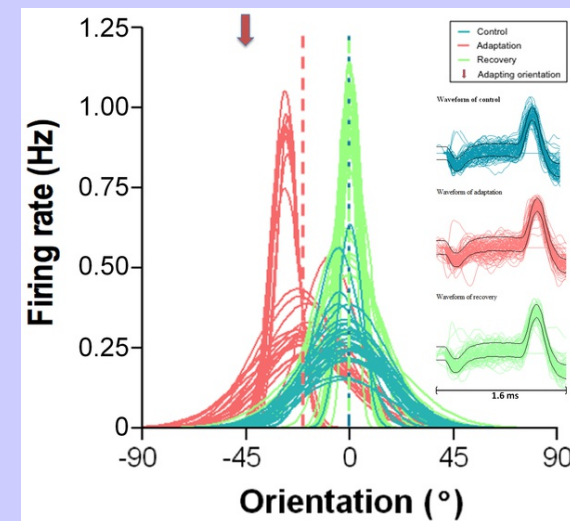




- But that's not how people normally plot it
- It might not make sense to sample *every possible* x
- Instead they might do some subset
 - For example, what if we just plot the points around the unit circle?



Tuning Curves



- Just along the unit circle



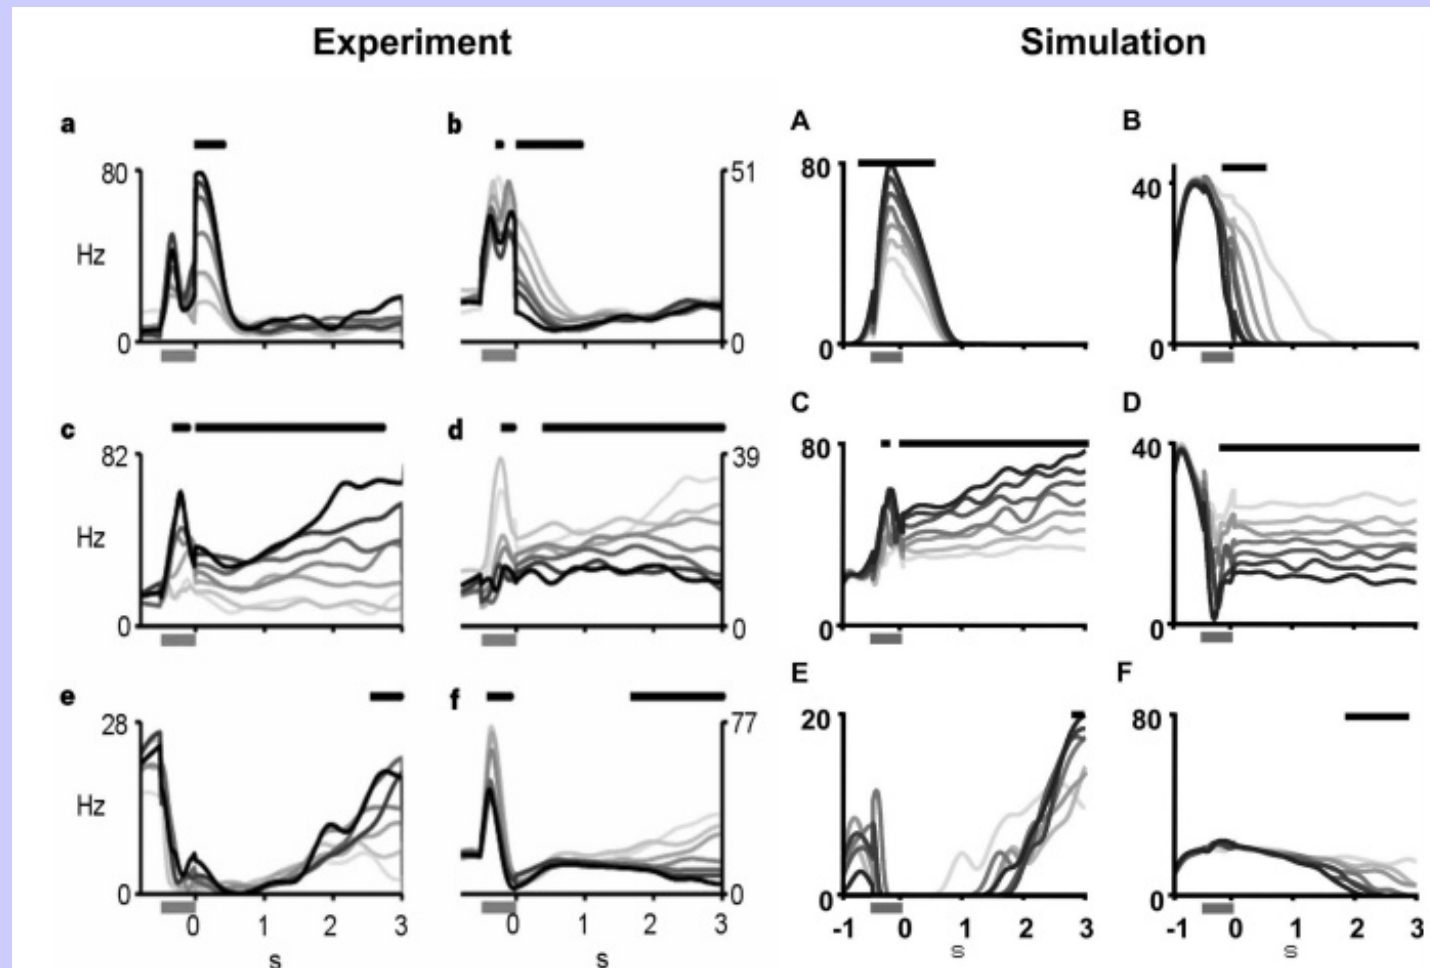
$$a_i = G_i [\alpha_i x \cdot e_i + J_i^{bias}]$$



Tuning Curves



(Georgopoulos et al., 1982)



(Singh & Eliasmith, 2005)



Tuning Curves

- General claim

- For any neural data, we can choose a space x such that we can match the neural data using

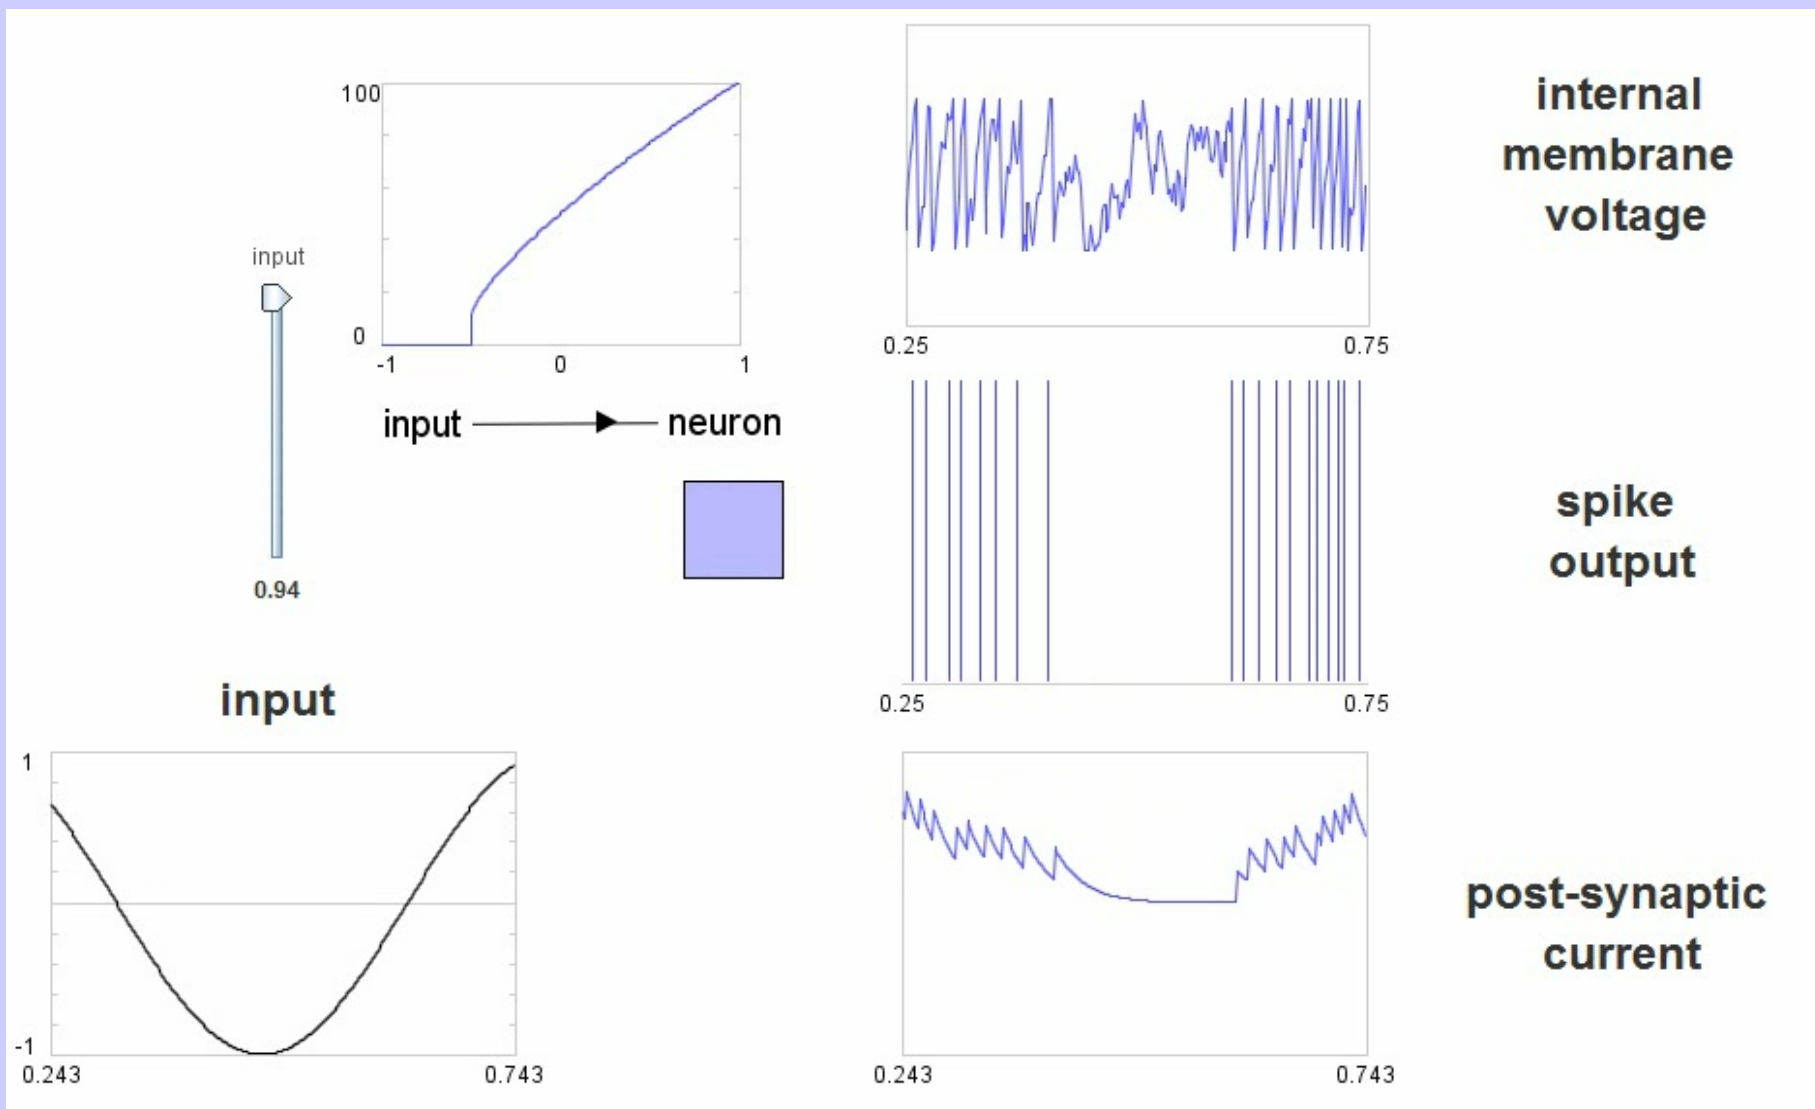
$$a_i = G_i[\alpha_i x \cdot e_i + J_i^{bias}]$$

- Note: we don't need this assumption for NEF to work. We just need some map from x to a .
 - But this form seems to work well
 - Many neurons respond to multiple things
 - And gives us a really interesting shortcut soon
- Note: what is e , physically?

- So what can we do with these neurons?

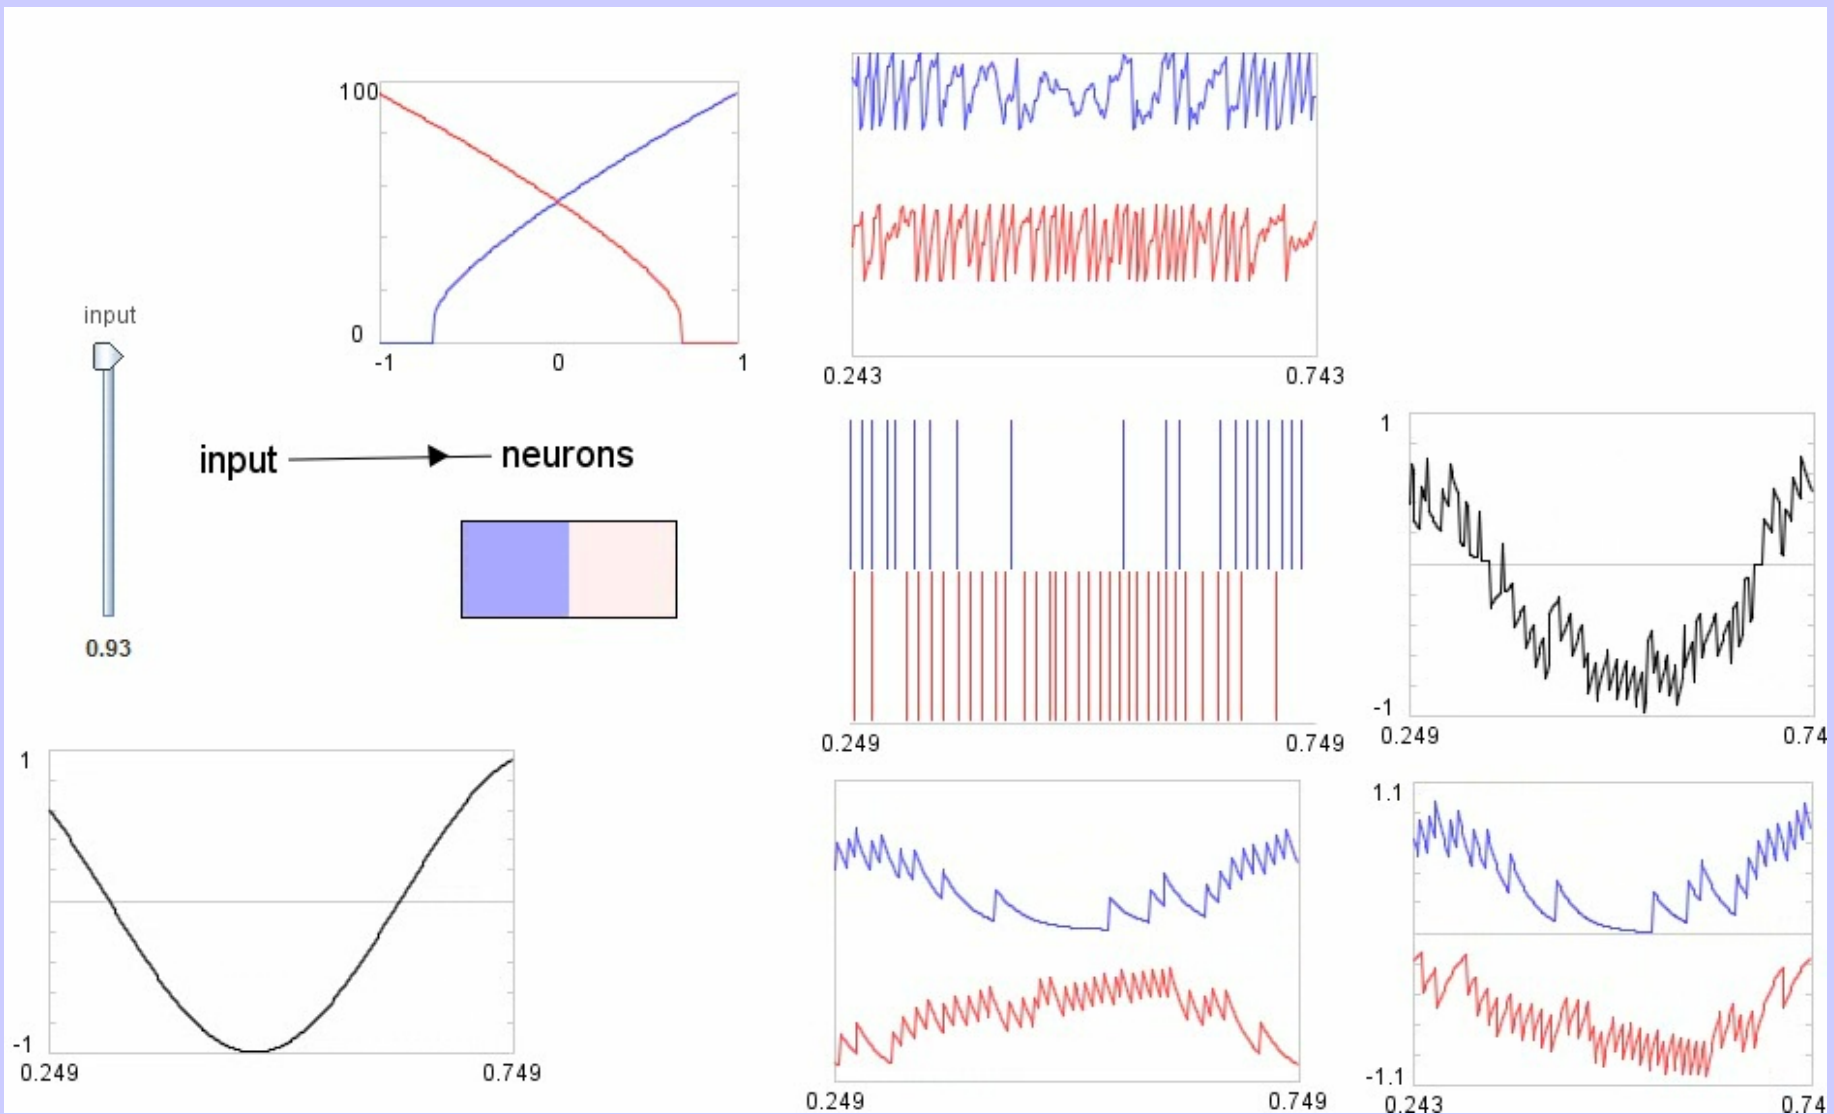


Single Neuron



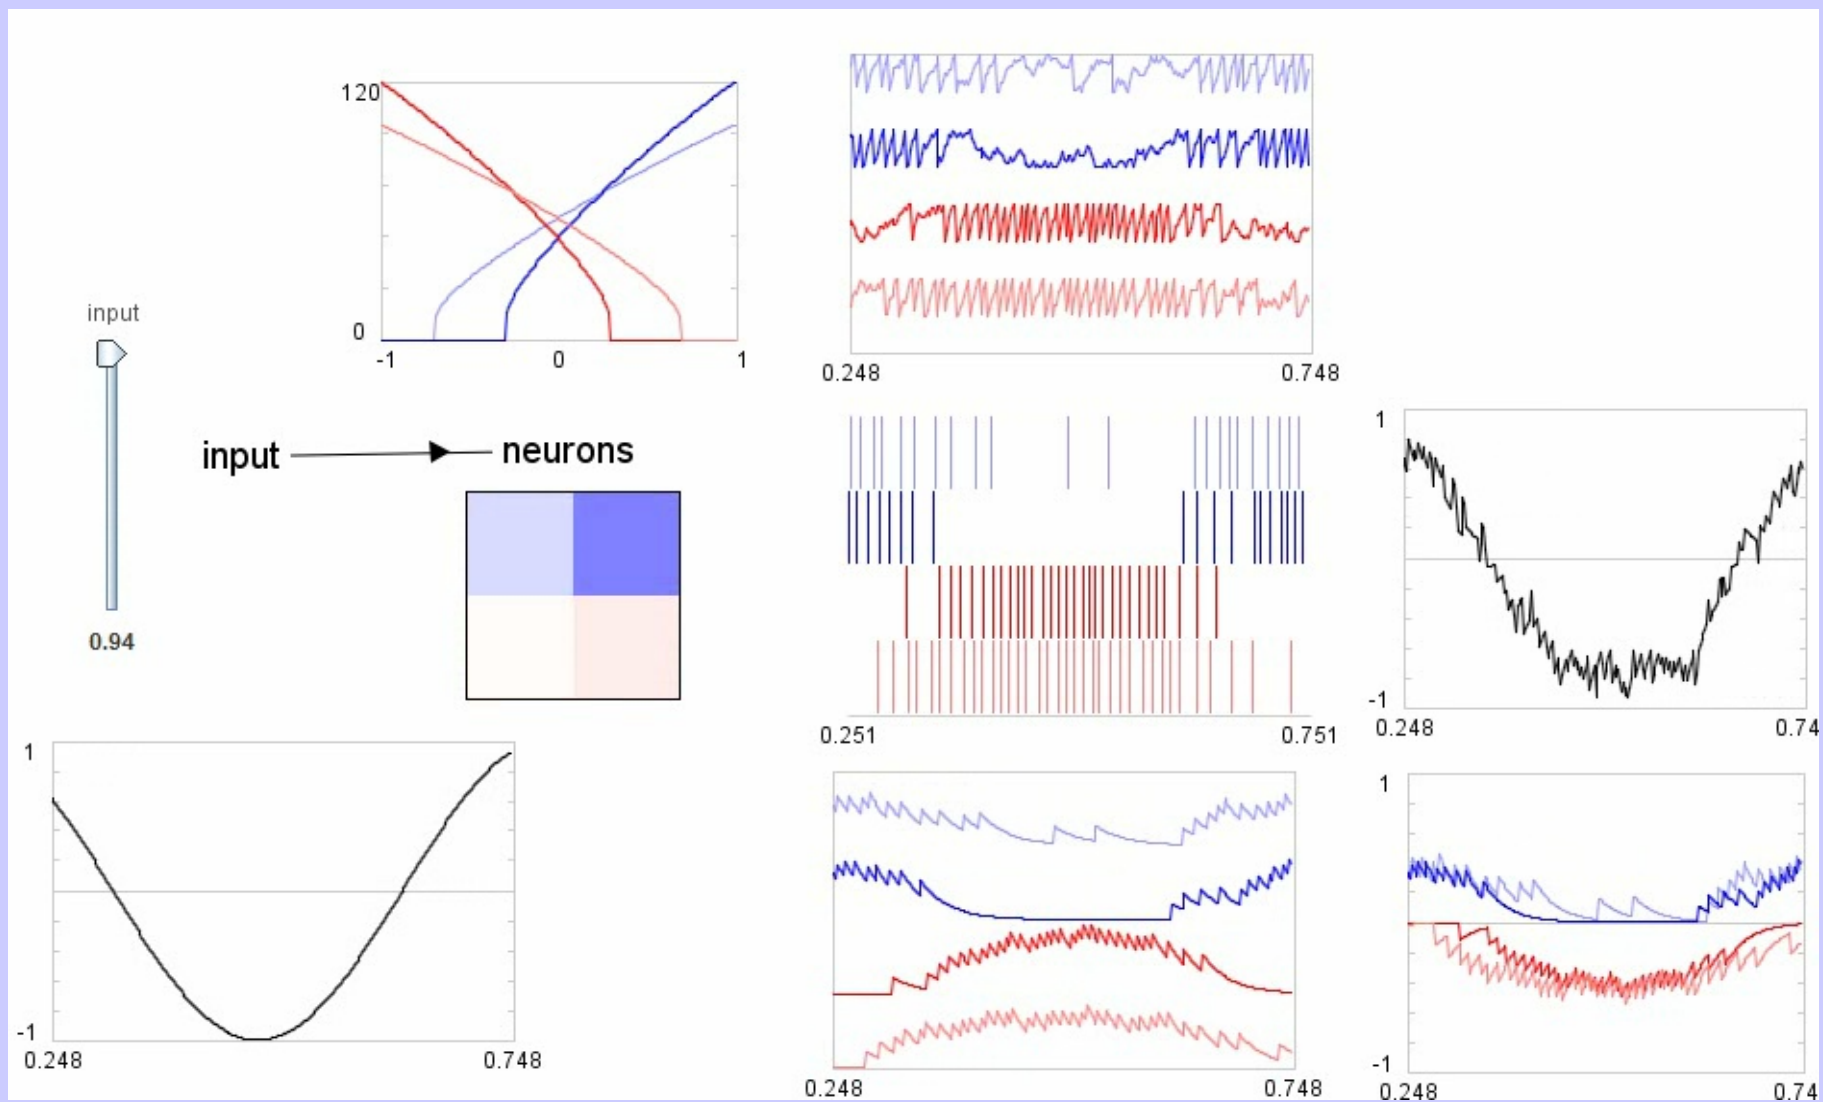


Two Neurons



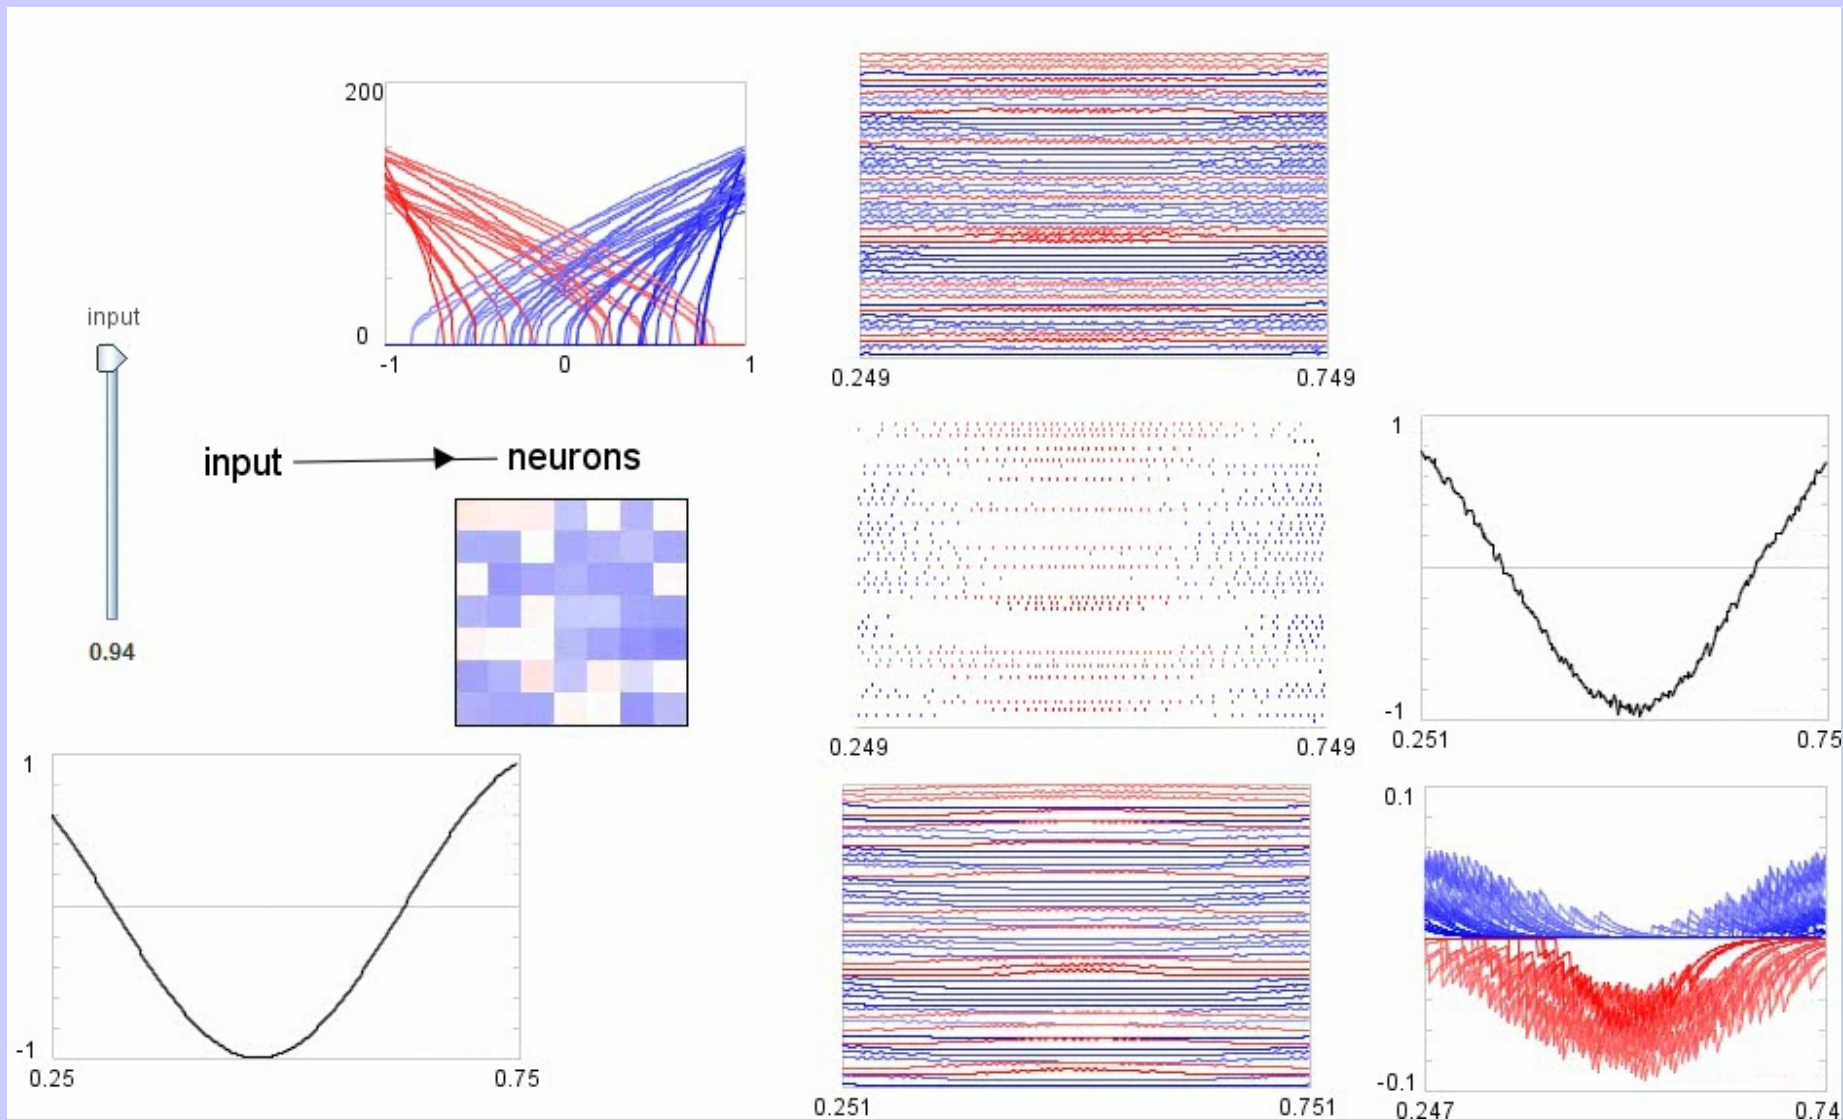


Four Neurons





Fifty Neurons





Neural computation

- Basic process

- A group of neurons stores some value \mathcal{X}
- Each neuron has some preferred stimulus e_i
- Current entering a neuron $J_i = x \cdot e_i$
- Neurons fire based on their input $a_i = G_i[J_i]$
- Decode output by weighted sum $\hat{x} = \sum_i a_i d_i$
- Find decoders by minimizing error $E = (x - \sum a_i d_i)^2$

- Extensions

- \mathcal{X} can be a scalar or a vector
- decoders for different functions $E = (f(x) - \sum a_i d_i)^2$




- [nengo example – scalar representation]



Multiple Dimensions

- Each neuron has a preferred direction (not just -1 or +1)
- Different weightings decode different values





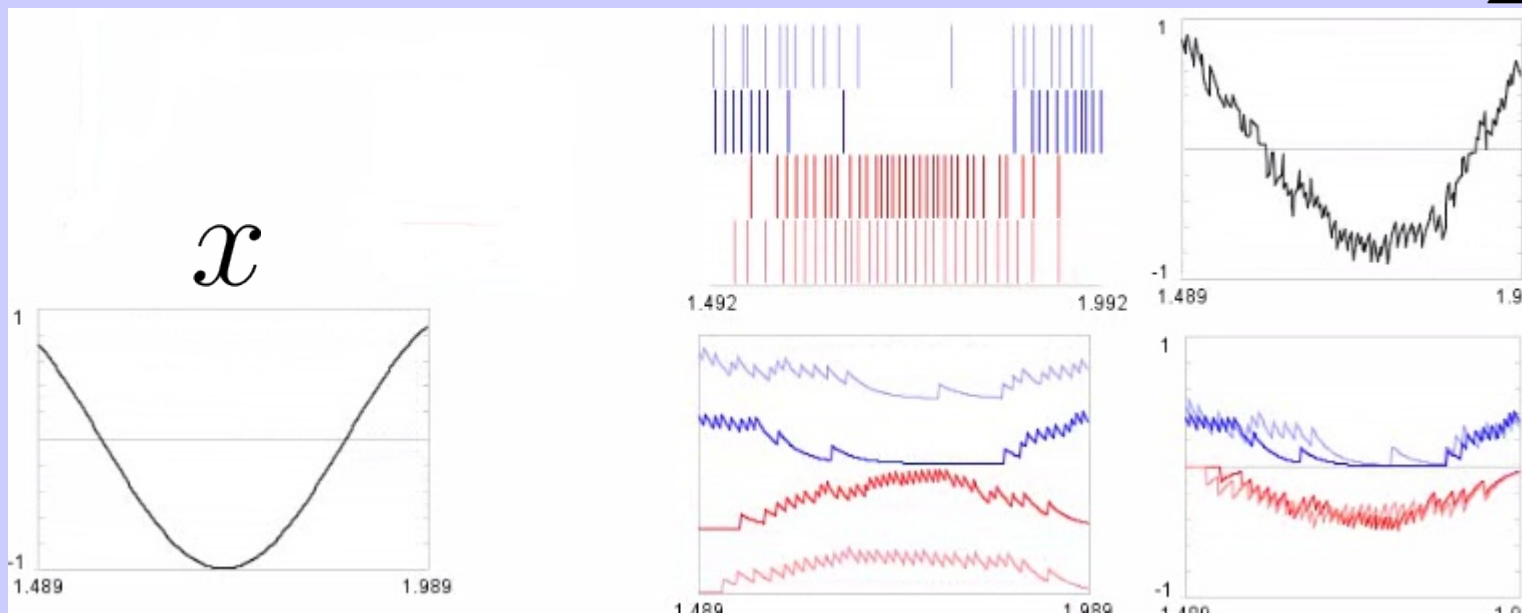
- [nengo example – 2d representation]



Decoders

- But where do we get d_i from?
 - $\hat{x} = \sum (a_i d_i)$
- Find the optimal d_i
 - How?

$$\hat{x} = \sum_i a_i d_i$$



$a_i d_i$



Decoders

- But where do we get d_i from?
 - $\hat{x} = \sum (a_i d_i)$
- Find the optimal d_i
 - How?

$$E = \frac{1}{2} \int_{-1}^1 (x - \sum_i (a_i d_i))^2 dx$$

- Take the derivative with respect to d_i

$$\frac{\partial E}{\partial d_i} = \frac{1}{2} \int_{-1}^1 2[x - \sum_j (a_j d_j)](-a_i) dx$$

$$\frac{\partial E}{\partial d_i} = - \int_{-1}^1 a_i x dx + \int_{-1}^1 \sum_j (a_j d_j a_i) dx$$

- At the minimum, $\frac{\partial E}{\partial d_i} = 0$

$$\int_{-1}^1 a_i x dx = \int_{-1}^1 \sum_j (a_j d_j a_i) dx$$

$$\int_{-1}^1 a_i x dx = \sum_j (\int_{-1}^1 a_i a_j dx) d_j$$



Decoders

- That's a system of N equations and N unknowns
- In fact, we can rewrite this in matrix form

$$\Upsilon = \Gamma d$$

where

$$\Upsilon_i = \frac{1}{2} \int_{-1}^1 a_i x dx$$

$$\Gamma_{ij} = \frac{1}{2} \int_{-1}^1 a_i a_j dx$$

- Do we have to do the integral over all x ?
 - Approximate the integral by sampling over x
 - S is the number of x values to use (S for samples)

$$\sum_x a_i x / S = \sum_j (\sum_x a_i a_j / S) d_j$$

$$\Upsilon = \Gamma d$$

where

$$\Upsilon_i = \sum_x a_i x / S$$

$$\Gamma_{ij} = \sum_x a_i a_j / S$$



- Notice that if A is the matrix of activities (the firing rate for each neuron for each x value), then $\Gamma = A^T A / S$ and $\Upsilon = A^T x / S$

So given

$$\Upsilon = \Gamma d$$

then


$$d = \Gamma^{-1} \Upsilon$$

or, equivalently

$$d_i = \sum_j \Gamma_{ij}^{-1} \Upsilon_j$$

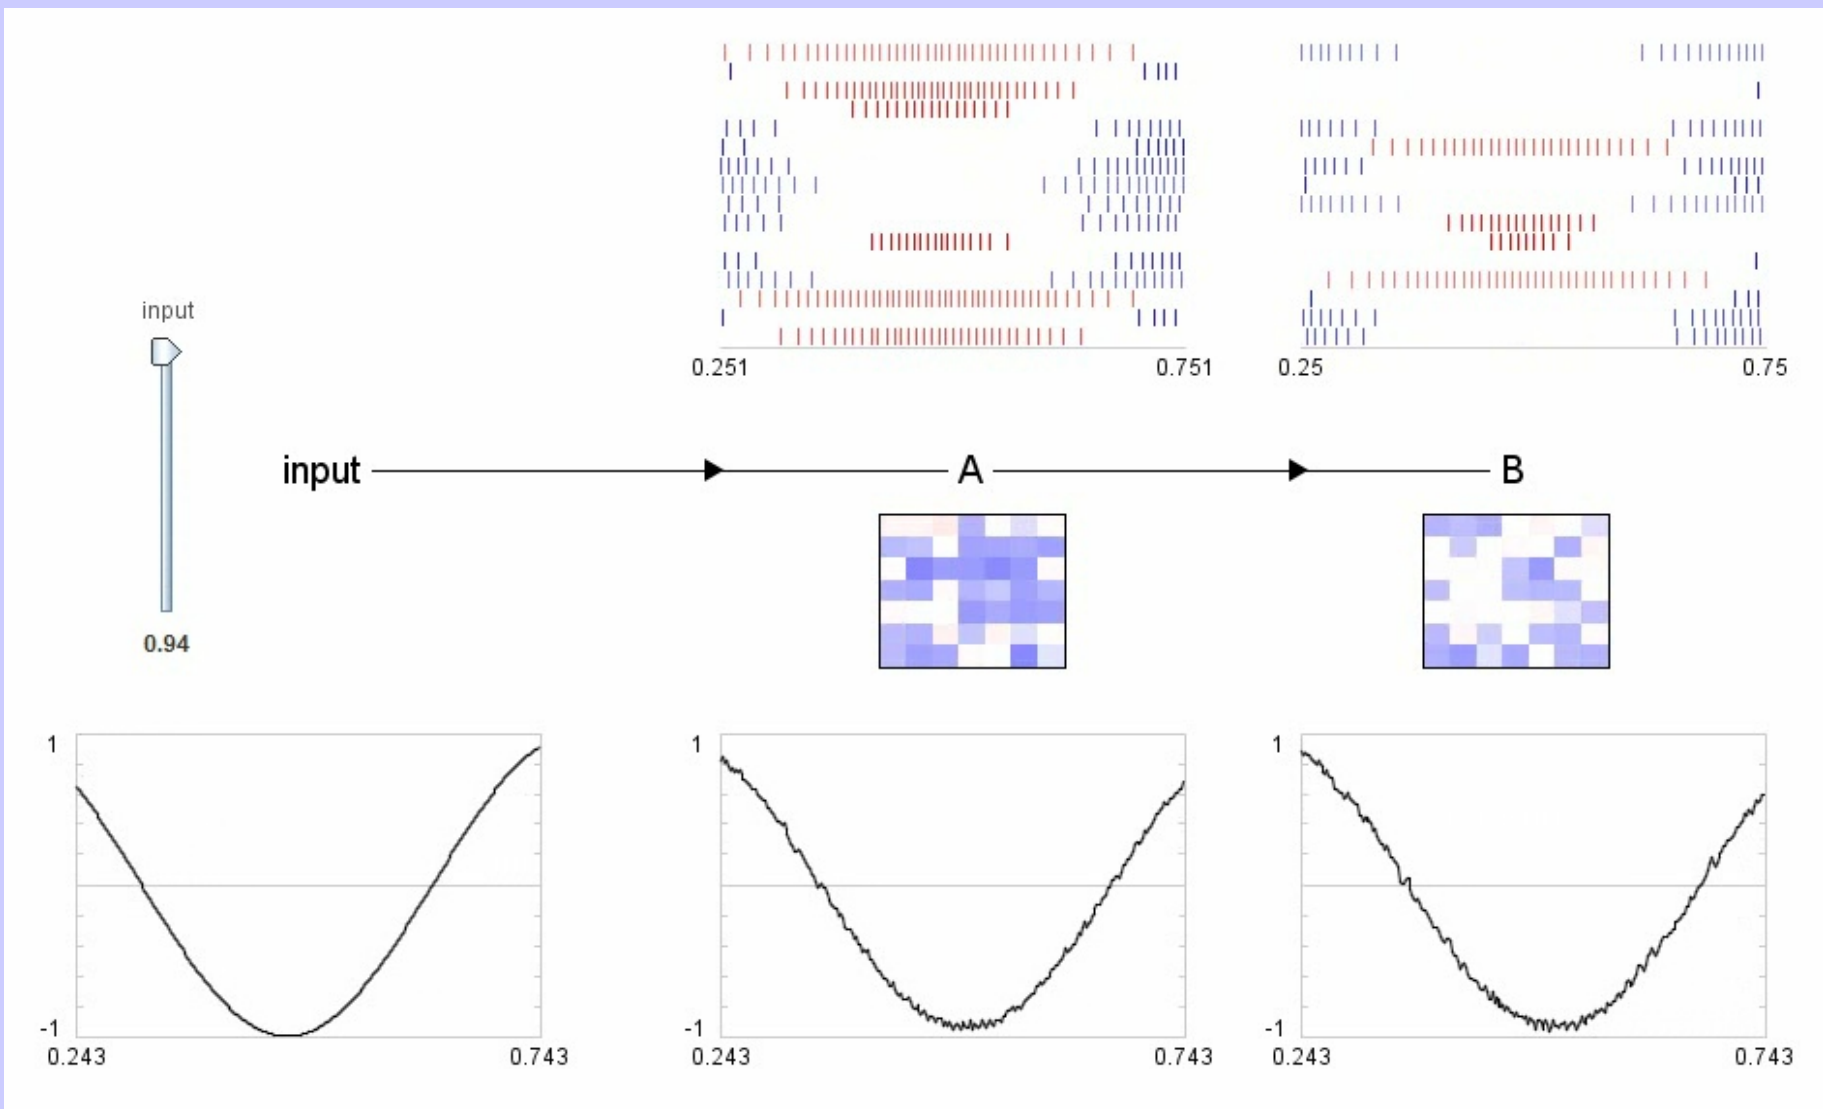


Connecting Neurons





Communication Channel

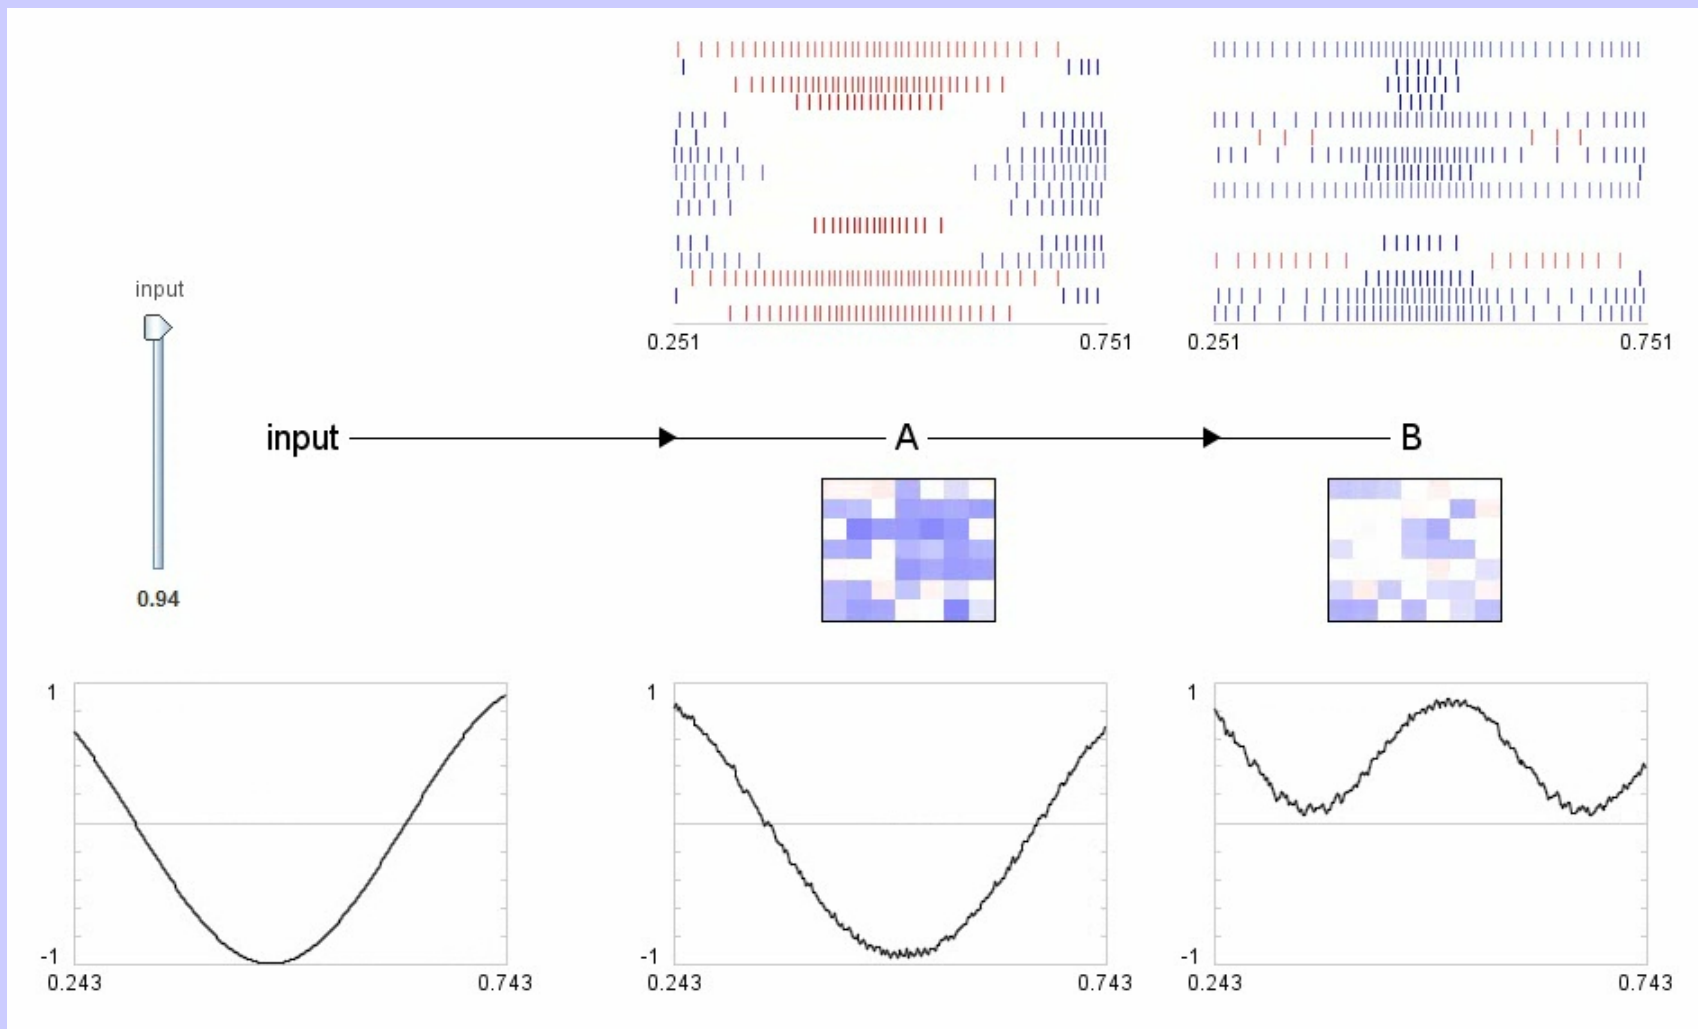




- [nengo example - communication]



Computation



Computing $y=x^2$ requires the same amount of effort as $y=x$



Neural computation

- [Nengo example: decoding functions]





Computation

- Addition?
 - [nengo example]
- Combination?
 - [nengo example]
- Multiplication?
 - [nengo example]



Neural Computation

- With enough neurons, we can approximate any function to any degree of accuracy
 - $MSE \propto 1/N$
- What functions are neurons good at approximating?
 - Do SVD on the tuning curves





- Low-degree polynomials (Legendre basis)



Computation

- Estimate any function $f(x)$
 - Accuracy increases as # neurons increases
 - Best at low-degree polynomials
- Not quite a perfect version of the function
 - Random noise due to neural activity
 - Smoothed due to post-synaptic current (varies from ~2ms to ~200ms)



$$f(x(t)) * h(t)$$



Biological Algorithms

- What do neural algorithms look like?
 - Each node (group of neurons) stores a vector
 - Each connection computes a function
 - and applies a filter
 - (set of functions and filter depends on neuron model)
- Different from standard connectionism
 - There, connections can only do linear weights
 - Some functions are easier than others
 - $\max(a,b)$ takes a very large number of neurons
 - $\sin(a+b) \cdot \cos(b - a)$ is pretty easy



Recurrent connections

- What happens if a group of neurons connects back to itself?
 - Depends on what function is being computed on the connection

$$f(x) = x + 1$$



$$f(x) = -x$$

$$f(x) = x^2$$



Nengo

- Open-source (free for non-commercial use)
 - <http://nengo.ca>
 - <http://github.com/nengo/nengo>
- Requirements
 - Python (2.7, 3.4, or 3.5)
 - NumPy
- Install
 - “pip install nengo”
 - “pip install nengo_gui”